

Restaurant revenue prediction

Team 38

Anjali Shenoy

Rehas Sachdeva

Saumya Rawat

Background

Enterprises today, highly motivated in the art of finding **anomalies, patterns and correlations** within data-sets.

Want to improve their **online reviews** to **attract clientele** or seek to establish a new business that is mindful of what **drives good reviews**, particularly true for **restaurants and food establishments**.



Introduction

Several studies conducted to look at the **correlation between a restaurant's success and its reviews and ratings.**

TFI is behind the famous brands like Burger King, Sbarro, Popeyes etc, interested in **extrapolating their data across geographies and cultures.**

We will be working with a **TFI data set of about 1 lakh Turkish restaurants.**

Problem Statement

Supervised learning problem

Objective

To develop a model and a set of preprocessing procedures to accurately predict the annual restaurant sales of 100,000 regional locations using various parameters.

Problem Challenges

- The size of training dataset is 137 samples while that of test dataset is 1,00,000 samples. This is a **large disparity**.
- **We don't have the ground truth for our test data.** So we cannot use sophisticated performance measures like precision, recall, k-fold cross validation etc. **We only know RMSE** for the entire test data.
- Whether to predict **revenue or log(revenue)** as we see that training data follows **normal distribution** when taken with log(revenue). But we can't say the same about test dataset.
- **Parsing** the data types of various attributes.
- **Unaccounted problem:** the disparity between the features for the training set and test set as the test set contains more information than the training set.
- **Categorical vs continuous problem:** whether the obfuscated P-Variables should be treated as categorical or continuous.
- **Zero problem:** For certain P-Variables, a large number of samples contain zero values and are dependent among each other such that if one p-variable has zero on a certain row, the probability that other p-variables take on a zero value is high.

Dataset Description

- The dataset based on 1 lakh Turkish restaurants, is uploaded on **Kaggle**.
- Size of training dataset: 137 samples, Size of test dataset: 1,00,000 samples.
- The **43 data fields** provided are:

ID: Restaurant ID

Open Date: Date that the restaurant opened in the format M/D/Y

City: The city name that the restaurant resides in

City Group: The type of city can be either big cities or other

Type: The type of the restaurant where FC - Food Court, IL - Inline, DT - Drive through and MB - Mobile.

P-Variables (P1, P2, ... ,P37): Obfuscated variables within three categories: demographic data, e.g population, age, gender; real estate data e.g car park availability and front facade; commercial data e.g points of interest, other vendors, etc. It is unknown if each variable contains a combination of the three categories or are mutually exclusive.

Revenue: Annual (transformed) revenue of a restaurant in a given year and is the target to be predicted.

Feature explanation

- Sales depend on the **location and type of city** it is in. If the city is a **metropolitan city**, it will have a **larger customer base** than a town, and hence revenue generated will be different in both these cities.
- On a similar note, revenues generated will be different for different restaurant types- **Drive through** will attract more customers in a remote area where as a **restaurant will attract more customers if it is placed in the heart of the city**.
- The open date attribute doesn't do much as such. But if processed to get **number of days a restaurant stays open, year of opening, month of opening**, we can get an idea of **cyclical or seasonal patterns** that affect revenue.
- Apart from these, we have **demographic attributes** e.g population, age, gender; **real estate data** e.g car park availability and front facade; and **commercial data** e.g points of interest, other vendors, etc. These also decide sales to varied extents. They can also be **correlated**.

Feature Extraction and Selection

- **Training and test dataset** - available as csv files containing 137 and 100,000 samples respectively. The input data is too large to be processed and majority of the data fields are **obfuscated variables** without giving any prior knowledge of each one. The data is pre processed by the following methods :
- We use **histograms to see that log(revenue)** follows an approximately **normal distribution**. So choosing target variable as log(revenue) instead of revenue improves performance of base models.
- **Opening date cannot simply be assumed to be a factor** so two additional features are created: month that they opened and the year that they opened. These two features can potentially help **proxy seasonality differences** since restaurant **revenues are highly cyclical**.
- Since the restaurant '**Type**'- '**MB**' **is not present in the training but available in the Test**, the Test set is modified so each mobile type restaurant is matched with a non-mobile type restaurant through finding the most similar features, as measured by euclidean distance.
- **Similarly for 'City' since the number of cities in the test set is more than the training set**, using KNN all 137 cities are clustered on the basis of P Variables that best describe Geographical locations and then these clusters are assigned to the city column of the datasets. To know exactly what each p-variable represents, under the assumption of mutually exclusive categories, a change in the mean over each city should elicit a change in certain p-variables. And using box plot of mean p-variables over each city, P1, P2, P11, P19, P20, P23, and P30 are identified to be approximately a good proxy for geographical location.
- We also use PCA for the P-Variables because we aren't given exactly what these represent. **They can be correlated**. So PCA can represent them better.

Validation techniques used

- We use **histograms to see that log(revenue)** follows an approximately **normal distribution**. So choosing target variable as log(revenue) instead of revenue improves performance of base models.
- **Visualizing a box plot of mean of P-variables** for each city helps us identify which P-variables are majorly geographical.
- **Davies-Bouldin index for K-Means clustering on variables:** P1,P2,P11,P19,P20,P23,P30 helps us validate the **best K** to be used as number of clusters.

Performance metrics

Root Mean Squared Error (RMSE)

Submissions are scored on the root mean squared error. RMSE is very common and is a suitable general-purpose error metric. Compared to the Mean Absolute Error, RMSE punishes large errors:

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2},$$

where \hat{y} is the predicted value and y is the original value.

Algorithms Mentioned in the paper

- **KNN** to account for **types of restaurants** that are there in test data but not in training data. The idea is to match those **unknown types with known types** (from training data) using nearest neighbour method based on rest of the attributes. **KNN is also used for zero problem**, to get an appropriate value for missing values of various attributes.
- **KMeans to account for cities** of restaurants that are there in test data but not in training data. The idea is to match those **unknown cities with known cities** (from training data) using clusters of cities.
- **PCA to reduce the dimensionality** of the data, especially useful for the p variables, because we are given no information about what they may represent, and so could themselves be **linear combinations** of the "real" variables, or simply be **highly dependent**.
- **Random forest, SVM: Support Vector Machines with Regression, and an ensemble method** based on combination of these two models, enable us to predict the annual revenue of a restaurant.

Analysis and Results

- Pre-processing on City and Type greatly improves performance over baseline models.
- Log transformation on revenue does not improve real test set performance although training set results are very promising.
- Treating zero problem with PCA or KNN doesn't improve performance probably due to large misspecification errors of the treatment models that introduce more noise rather than clarity.
- Apart from zero problem treatment, all solutions proposed in the paper together lead to least RMSE.

TODO for final evaluation

- Implement inbuilt functions used in phase 1.
- Improve analysis of seasonal revenue, if possible to get relevant data.
- Implement neural network based model.
- Compare various models and visualize results.

Languages and Toolkits

- Python, sklearn
- MATLAB
- SAS
- JMP
- t-SNE

Scope

- The solution is **only applicable to Turkish restaurants** and locations on which the data is based. A different location based data may need a different kind of ensemble for accuracy.
- It is limited to only **annual and not seasonal** revenue analysis.
- Our scope is also **limited by the neural network model** that we build.

Timeline

